

POINTER ANALYSIS

PROGRAM ANALYSIS AND OPTIMIZATION – DCC888

Fernando Magno Quintão Pereira

fernando@dcc.ufmg.br

Pointers

- Pointers are a feature of imperative programming languages.
- They are very useful, because they avoid the need to copy entire data-structures when passing information from one program routine to another.
- Nevertheless, pointers make it very hard to reason about programs.
- The direct consequence of these difficulties is that compilers have a hard time trying to understand and modify imperative programs.
 - Compilers's main ally in this battle is the *pointer analysis*, also called alias analysis, or points-to analysis.

The Need for Pointer Analysis

```
#include <stdio.h>
int main() {
    int i = 7;
    int* p = &i;
    *p = 13;
    printf("The value of i = %d\n", i);
}
```

- 1) What does the program on the left do?
- 2) How could this program be optimized?
- 3) Which information is needed to optimize this program?
- 4) Do you think gcc -O1 can optimize this program?

The Need for Pointer Analysis

```
#include <stdio.h>
int main() {
    int i = 7;
    int* p = &i;
    *p = 13;
    printf("The value of i = %d\n", i);
}
```

```
$> gcc -O1 pointer.c -S
```

Which information
was needed to
optimize the target
program?

```
_main:
    pushl    %ebp
    movl     %esp, %ebp
    pushl    %ebx
    subl     $20, %esp
    call     L3
"L0000000000001$pb":
L3:
    popl     %ebx
    movl     $13, 4(%esp)
    leal     LC0-"L1$pb"(%ebx), %eax
    movl     %eax, (%esp)
    call     _printf
    addl     $20, %esp
    popl     %ebx
    leave
    ret
.subsections_via_symbols
```

The Need for Pointer Analysis

```
void sum0(int* a, int* b, int* r, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        r[i] = a[i];  
        if (!b[i]) {  
            r[i] = b[i];  
        }  
    }  
}
```

How could we
optimize this
program?

The Need for Pointer Analysis

```
void sum0(int* a, int* b, int* r, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        r[i] = a[i];  
        if (!b[i]) {  
            r[i] = b[i];  
        }  
    }  
}
```

```
void sum1(int* a, int* b, int* r, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        int tmp = a[i];  
        if (!b[i]) {  
            tmp = b[i];  
        }  
        r[i] = tmp;  
    }  
}
```

How much
faster do you
think is sum1?

The Need for Pointer Analysis

```
$> time ./a.out  
sum0
```

0	1	0	3
0	11	0	13

```
real 0m6.299s  
user 0m6.285s  
sys 0m0.008s
```

```
$> time ./a.out a  
sum1
```

0	1	0	3
0	11	0	13

```
real 0m1.345s  
user 0m1.340s  
sys 0m0.004s
```

```
int main(int argc, char** argv) {  
    int* a = (int*) malloc(SIZE * 4);  
    int* b = (int*) malloc(SIZE * 4);  
    int* c = (int*) malloc(SIZE * 4);  
    int i;  
    for (i = 0; i < SIZE; i++) {  
        a[i] = i;  
        b[i] = i%2;  
    }  
    if (argc % 2) {  
        printf("sum0\n");  
        for (i = 0; i < LOOP; i++) {  
            sum0(a, b, c, SIZE);  
        }  
    } else {  
        printf("sum1\n");  
        for (i = 0; i < LOOP; i++) {  
            sum1(a, b, c, SIZE);  
        }  
    }  
    print(c, 20);  
}
```

If you want to test it...

```
#include <stdio.h>
#include <stdlib.h>
void sum0(int* a, int* b, int* r, int N) {
    int i;
    for (i = 0; i < N; i++) {
        r[i] = a[i];
        if (!b[i]) {
            r[i] = b[i];
        }
    }
}
void sum1(int* a, int* b, int* r, int N) {
    int i;
    for (i = 0; i < N; i++) {
        int tmp = a[i];
        if (!b[i]) {
            tmp = b[i];
        }
        r[i] = tmp;
    }
}
void print(int* a, int N) {
    int i;
    for (i = 0; i < N; i++) {
        if (i % 10 == 0) {
            printf("\n");
        }
        printf("%8d", a[i]);
    }
}
```

```
#define SIZE 10000
#define LOOP 100000

int main(int argc, char** argv) {
    int* a = (int*) malloc(SIZE * 4);
    int* b = (int*) malloc(SIZE * 4);
    int* c = (int*) malloc(SIZE * 4);
    int i;
    for (i = 0; i < SIZE; i++) {
        a[i] = i;
        b[i] = i%2;
    }
    if (argc % 2) {
        printf("sum0\n");
        for (i = 0; i < LOOP; i++) {
            sum0(a, b, c, SIZE);
        }
    } else {
        printf("sum1\n");
        for (i = 0; i < LOOP; i++) {
            sum1(a, b, c, SIZE);
        }
    }
    print(c, 20);
}
```


The Need for Pointer Analysis

```
void sum0(int* a, int* b, int* r, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        r[i] = a[i];  
        if (!b[i]) {  
            r[i] = b[i];  
        }  
    }  
}
```

```
void sum1(int* a, int* b, int* r, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        int tmp = a[i];  
        if (!b[i]) {  
            tmp = b[i];  
        }  
        r[i] = tmp;  
    }  
}
```

This code **here**
is quite smart.
Do you
understand it?

```
L4:  
    movl    (%edi,%edx,4), %eax  
    movl    %eax, (%ecx,%edx,4)  
    cmpl    $0, (%esi,%edx,4)  
    jne     L5  
    movl    $0, (%ecx,%edx,4)  
L5:  
    incl    %edx  
    cmpl    %ebx, %edx  
    jne     L4
```

```
L12:  
    cmpl    $0, (%esi,%edx,4)  
    movl    $0, %eax  
    cmovne  (%edi,%edx,4), %eax  
    movl    %eax, (%ebx,%edx,4)  
    incl    %edx  
    cmpl    %ecx, %edx  
    jne     L12
```

The Need for Pointer Analysis

```
void sum0(int* a, int* b, int* r, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        r[i] = a[i];  
        if (!b[i]) {  
            r[i] = b[i];  
        }  
    }  
}
```

Why is gcc -O1
unable to
optimize this
program?

The Need for Pointer Analysis

```
void sum0(int* a, int* b, int* r, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        r[i] = a[i];  
        if (!b[i]) {  
            r[i] = b[i];  
        }  
    }  
}
```

Would both programs produce the same result if the second and third parameters were pointers to the same array?

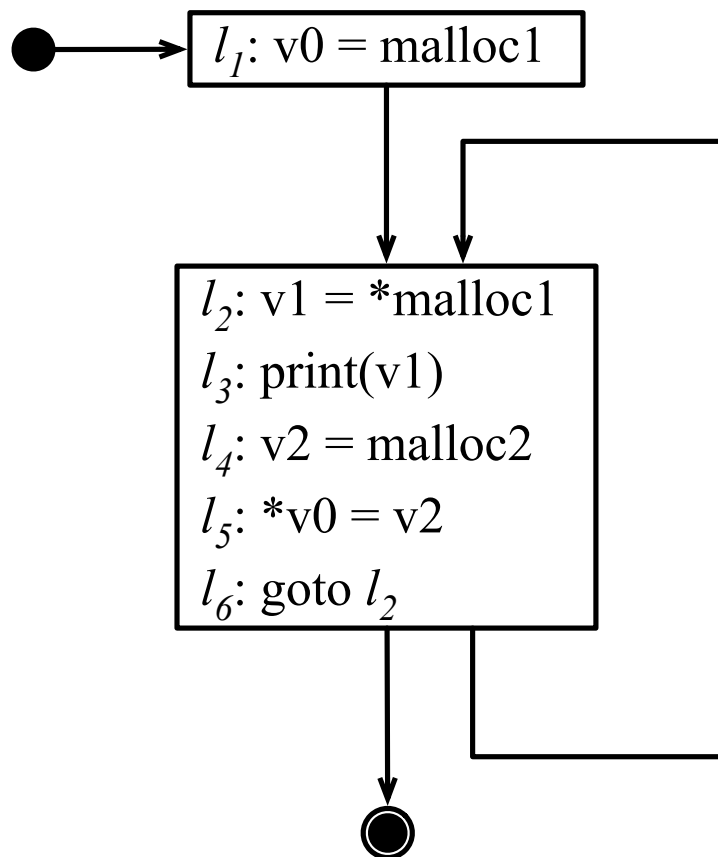
```
void sum1(int* a, int* b, int* r, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        int tmp = a[i];  
        if (!b[i]) {  
            tmp = b[i];  
        }  
        r[i] = tmp;  
    }  
}
```

Pointer Analysis

- The goal of pointer analysis is to determine which are the memory locations pointed by each pointer in the program.
- Pointer analysis is usually described and solved as a constraint based analysis.
- The most efficient algorithm that we know about is $O(n^3)$.
 - This complexity is very high, and pointer analysis takes too long in very large programs.
- After register allocation, pointer analysis is possibly the most studied topic inside the science of compiler design.
 - Research aims at speed and precision.

Can we solve Pointer Analysis with Data-flow?

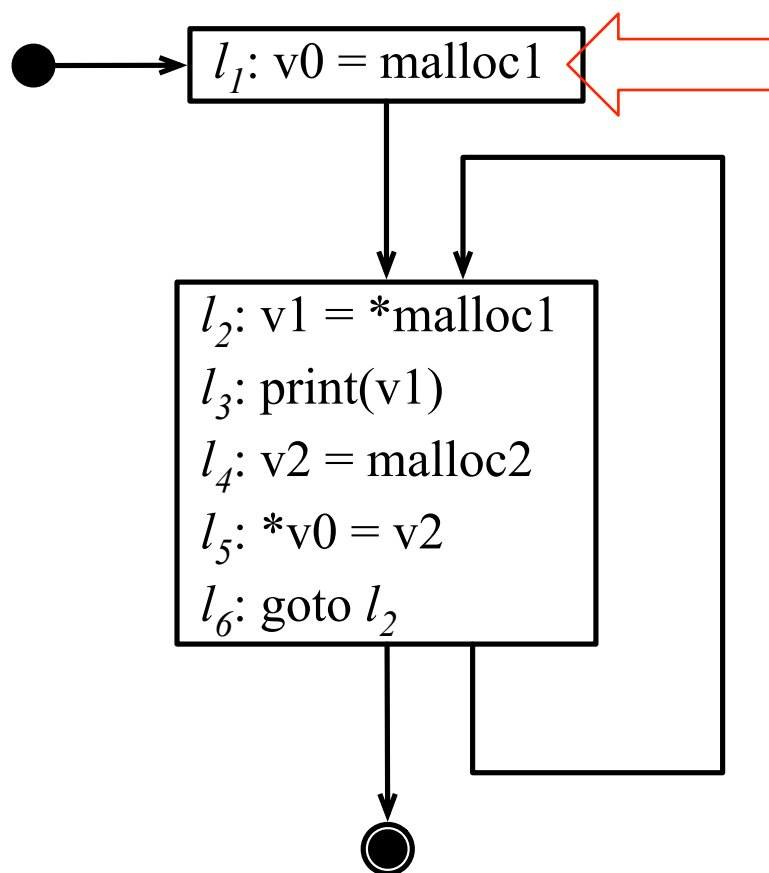
- We must determine, for each variable in the program, the set of memory locations that the variable may point to.



What are the points-to sets associated with the variables $v0$, $v1$, $v2$, malloc1 and malloc2 in this program?

Can we solve Pointer Analysis with Data-flow?

- We must determine, for each variable in the program, the set of memory locations that the variable may point to.

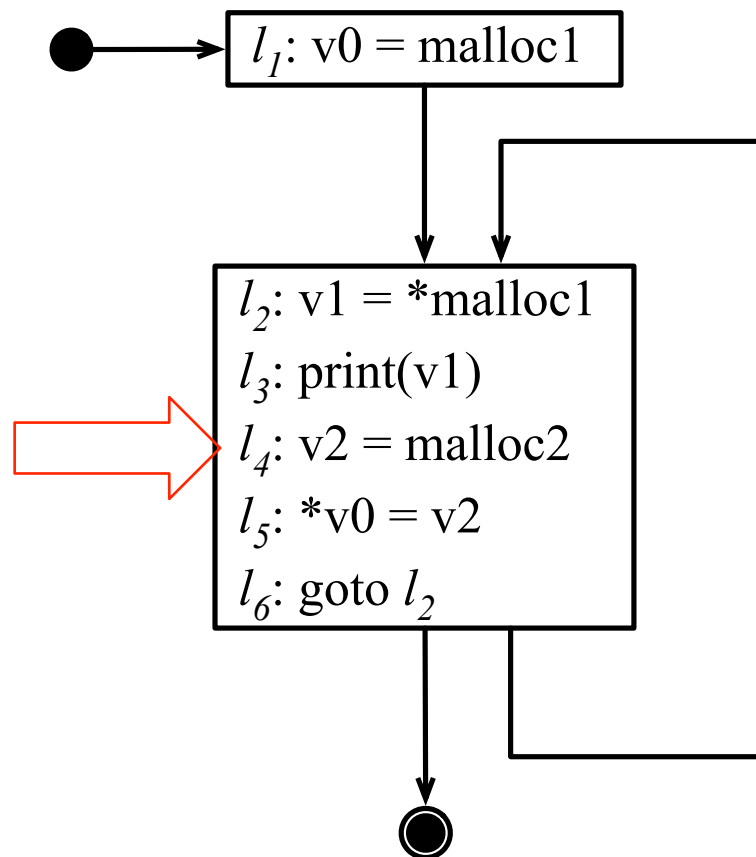


First of all, we know that $\text{malloc1} \in P(v0)$, due to the assignment at label 1

$$P(v0) \supseteq \{\text{malloc1}\}$$

Can we solve Pointer Analysis with Data-flow?

- We must determine, for each variable in the program, the set of memory locations that the variable may point to.

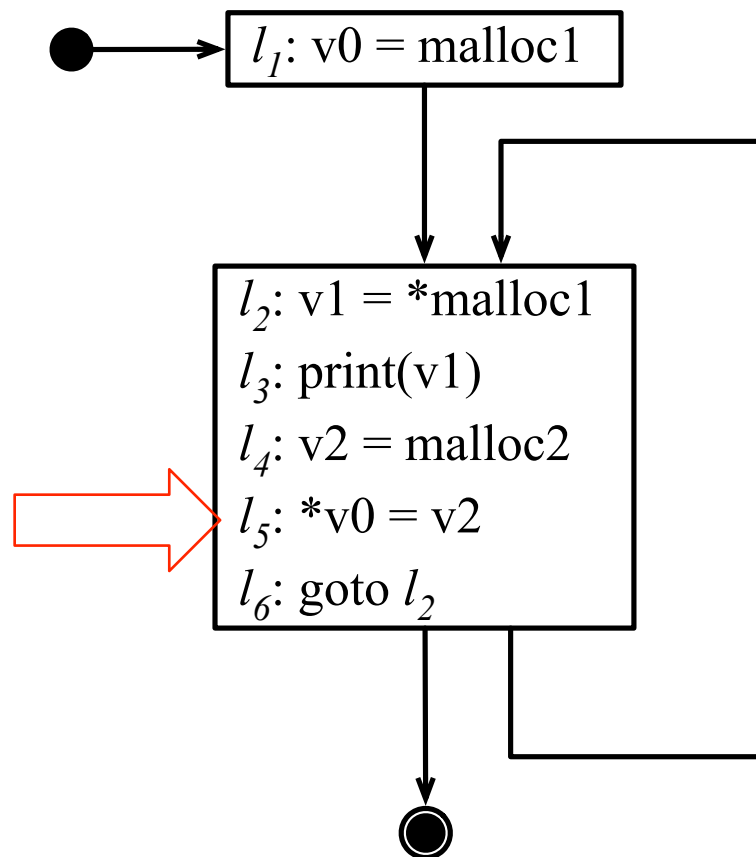


We also know that
 $\text{malloc2} \in P(v2)$, due
to the assignment at
label 4

$P(v0) \supseteq \{\text{malloc1}\}$
 $P(v2) \supseteq \{\text{malloc2}\}$

Can we solve Pointer Analysis with Data-flow?

- We must determine, for each variable in the program, the set of memory locations that the variable may point to.



Line 5 is weird: it adds to the points-to set of the things pointed by $v0$ the points-to set of $v2$. At this points, $P(v0) = \{\text{malloc1}\}$

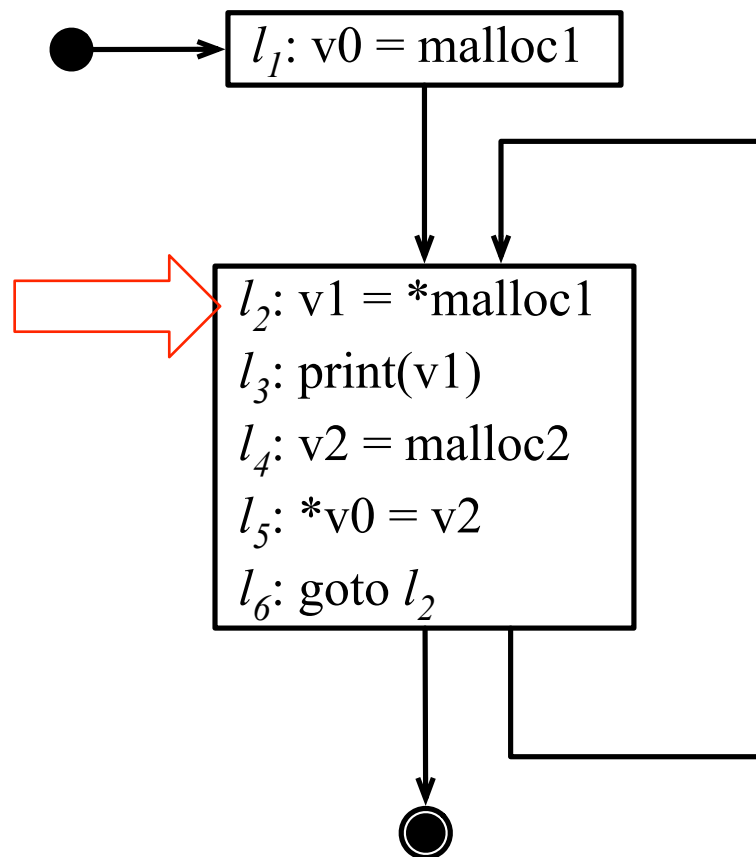
$P(v0) \supseteq \{\text{malloc1}\}$

$P(v2) \supseteq \{\text{malloc2}\}$

$P(\text{malloc1}) \supseteq \{\text{malloc2}\}$

Can we solve Pointer Analysis with Data-flow?

- We must determine, for each variable in the program, the set of memory locations that the variable may point to.

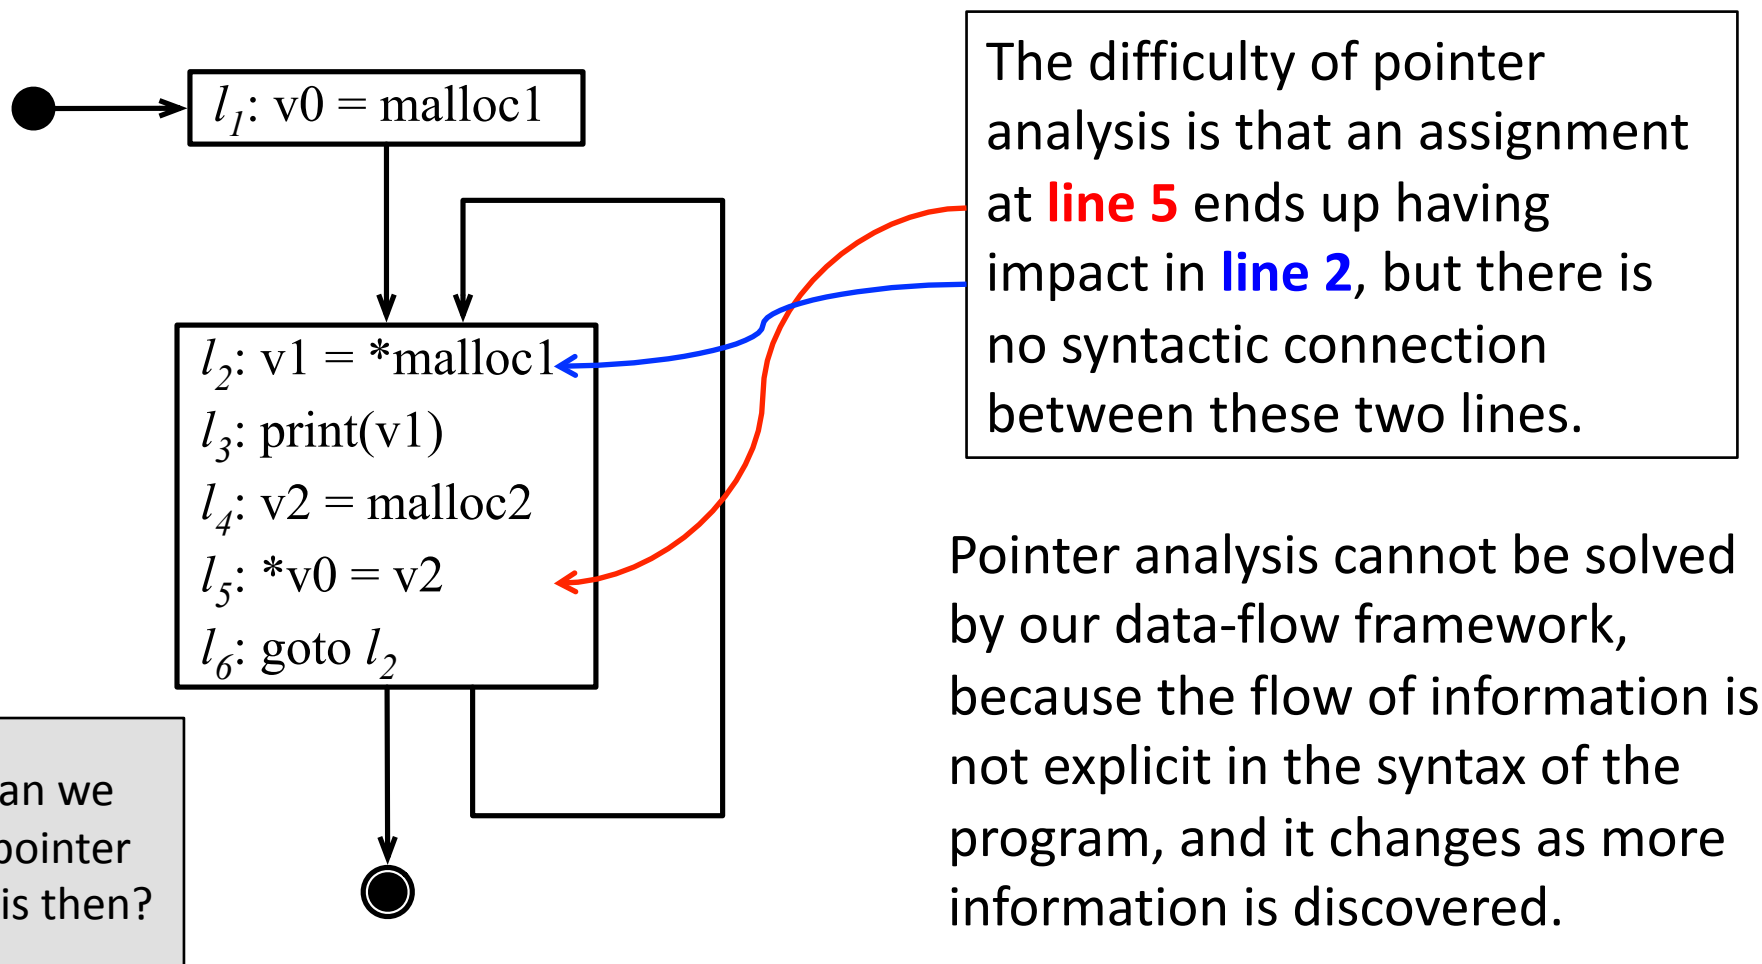


But, now $P(\text{malloc1})$ has been changed, and line 2 start being important: $P(v1)$ must include the points-to sets of things pointed by malloc1

$P(v0) \supseteq \{\text{malloc1}\}$
 $P(v2) \supseteq \{\text{malloc2}\}$
 $P(\text{malloc1}) \supseteq \{\text{malloc2}\}$
 $P(v1) \supseteq \{\text{malloc2}\}$

Can we solve Pointer Analysis with Data-flow?

- We must determine, for each variable in the program, the set of memory locations that the variable may point to.



Fernando Magno Quintão Pereira

`fernando@dcc.ufmg.br`

`lac.dcc.ufmg.br`

ANDERSEN'S POINTER ANALYSIS



Solving Pointer Analysis with Constraints

- Four different constructs typically found in imperative programming languages give us constraints:

Statement	Constraint name	Constraint
$a = \&b$	base	$P(a) \supseteq \{b\}$
$a = b$	simple	$P(a) \supseteq P(b)$
$a = *b$	load	$t \in P(b) \Rightarrow P(a) \supseteq P(t)$
$*a = b$	store	$t \in P(a) \Rightarrow P(t) \supseteq P(b)$

Are these constraints similar to those seen in the control flow analysis?

- This type of constraints are *inclusion based*. This pointer analysis is called *Andersen style*, after Lars Andersen, who first described the inclusion based points-to analysis[♠].

[♠]: Program Analysis and Specialization for the C Programming Language, 1994

Solving Points-to Analysis with Graphs

- We can use an approach similar to that seen in the constraint based analyses to solve points-to analysis; however, notice that the constraints are not exactly the same:

Constraint of Pointer Analysis

$$P(a) \supseteq \{b\}$$

$$P(a) \supseteq P(b)$$

$$t \in P(b) \Rightarrow P(a) \supseteq P(t)$$

$$t \in P(a) \Rightarrow P(t) \supseteq P(b)$$

Constraint of Control Flow Analysis

$$lhs \subseteq rhs$$

$$\{t\} \subseteq rhs' \Rightarrow lhs \subseteq rhs$$

Do you remember
the graph based
algorithm to solve
constraint based
analyses?

Solving Points-to Analysis with Graphs

- We will solve points-to analysis with the same algorithm that we have used to solve control flow analysis.
- We will show several improvements on that algorithm that we can use to make it scale to handle very large points-to sets.

Constraint of Pointer Analysis

$$P(a) \supseteq \{b\}$$

$$P(a) \supseteq P(b)$$

$$t \in P(b) \Rightarrow P(a) \supseteq P(t)$$

$$t \in P(a) \Rightarrow P(t) \supseteq P(b)$$

Constraint of Control Flow Analysis

$$lhs \subseteq rhs$$

$$\{t\} \subseteq rhs' \Rightarrow lhs \subseteq rhs$$

The Points-to Graph

- The points-to graph is a graph (V, E) used to solve the pointer analysis.
 - The graph has a node for each variable v in the constraint system.
 - Each node is associated with a points-to set $P(v)$
 - The graph has an edge (v_1, v_2) if $v_1 \subseteq v_2$
 - Initially, the points-to graph has an edge for each constraint $P(v_1) \supseteq P(v_2)$ in the constraint system.

`b = &a`

`a = &c`

`d = a`

`*d = b`

`a = *d`

How are the
constraints for
this program?

The Points-to Graph

- The points-to graph is a graph (V, E) used to solve the pointer analysis.
 - The graph has a node for each variable v in the constraint system.
 - Each node is associated with a points-to set $P(v)$
 - The graph has an edge (v_1, v_2) if $v_1 \subseteq v_2$
 - Initially, the points-to graph has an edge for each constraint $P(v_1) \supseteq P(v_2)$ in the constraint system.

$b = \&a$

$a = \&c$

$d = a$

$*d = b$

$a = *d$

How is the initial
points-to graph
for the program
on the left?

$P(b) \supseteq \{a\}$

$P(a) \supseteq \{c\}$

$P(d) \supseteq P(a)$

$t \in P(d) \Rightarrow P(t) \supseteq P(b)$

$t \in P(d) \Rightarrow P(a) \supseteq P(t)$

The Points-to Graph

- The points-to graph is a graph (V, E) used to solve the pointer analysis.
 - The graph has a node for each variable v in the constraint system.
 - Each node is associated with a points-to set $P(v)$
 - The graph has an edge (v_1, v_2) if $v_1 \subseteq v_2$
 - Initially, the points-to graph has an edge for each constraint $P(v_1) \supseteq P(v_2)$ in the constraint system.

$$P(b) \supseteq \{a\}$$

$$P(a) \supseteq \{c\}$$

$$P(d) \supseteq P(a)$$

$$t \in P(d) \Rightarrow P(t) \supseteq P(b)$$

$$t \in P(d) \Rightarrow P(a) \supseteq P(t)$$

How can we
solve the rest of
the points-to
analysis?



The Iterative Solver

- Once we have an initial points-to graph, we iterate over the load and store constraints (henceforth called complex constraints), alternating two steps:
 - Adding new edges to the graph
 - Propagating points-to information
- This algorithm effectively builds the transitive closure of the points-to graph


```
let  $G = (V, E)$   
 $W = V$   
while  $W \neq []$  do  
   $n = \text{hd}(W)$ 
```

```
    for each  $v \in P(n)$  do  
      for each load " $a = *n$ " do  
        if  $(v, a) \notin E$  then  
           $E = E \cup \{(v, a)\}$   
           $W = v :: W$   
      for each store " $*n = b$ " do  
        if  $(b, v) \notin E$  then  
           $E = E \cup \{(b, v)\}$   
           $W = b :: W$ 
```

```
    for each  $(n, z) \in E$  do  
       $P(z) = P(z) \cup P(n)$   
      if  $P(z)$  has changed then  
         $W = z :: W$ 
```

The Iterative Solver

How the graph will look like after the initial propagation phase?




```

let  $G = (V, E)$ 
 $W = V$ 
while  $W \neq []$  do
   $n = \text{hd}(W)$ 
  for each  $v \in P(n)$  do
    for each load " $a = *n$ " do
      if  $(v, a) \notin E$  then
         $E = E \cup \{(v, a)\}$ 
         $W = v::W$ 
    for each store " $*n = b$ " do
      if  $(b, v) \notin E$  then
         $E = E \cup \{(b, v)\}$ 
         $W = b::W$ 
  for each  $(n, z) \in E$  do
     $P(z) = P(z) \cup P(n)$ 
    if  $P(z)$  has changed then
       $W = z::W$ 
  
```

The Iterative Solver

*d = b
a = *d

Now, we have two complex constraints to evaluate. How will be the first evaluation?



```

let G = (V, E)
W = V
while W ≠ [] do
  n = hd(W)
  for each v ∈ P(n) do
    for each load "a = *n" do
      if (v, a) ∉ E then
        E = E ∪ {(v, a)}
        W = v::W
    for each store "*n = b" do
      if (b, v) ∉ E then
        E = E ∪ {(b, v)}
        W = b::W
  for each (n, z) ∈ E do
    P(z) = P(z) ∪ P(n)
    if P(z) has changed then
      W = z::W
  
```


The Iterative Solver

*d = b

a = *d



How will be the graph
after we evaluate the
load?



let $G = (V, E)$

$W = V$

while $W \neq []$ **do**

$n = \text{hd}(W)$

for each $v \in P(n)$ **do**

→ **for each** load " $a = *n$ " **do**

if $(v, a) \notin E$ **then**

$E = E \cup \{(v, a)\}$

$W = v::W$

for each store " $*n = b$ " **do**

if $(b, v) \notin E$ **then**

$E = E \cup \{(b, v)\}$

$W = b::W$

for each $(n, z) \in E$ **do**

$P(z) = P(z) \cup P(n)$


if $P(z)$ has changed **then**

$W = z::W$

The Iterative Solver

*d = b
a = *d

Are we done, or has any
points-to set changed;
hence, forcing a new
iteration?



```

let G = (V, E)
W = V
while W ≠ [] do
  n = hd(W)
  for each v ∈ P(n) do
    for each load "a = *n" do
      if (v, a) ∉ E then
        E = E ∪ {(v, a)}
        W = v::W
    for each store "*n = b" do
      if (b, v) ∉ E then
        E = E ∪ {(b, v)}
        W = b::W
  for each (n, z) ∈ E do
    P(z) = P(z) ∪ P(n)
    if P(z) has changed then
      W = z::W
  
```


The Iterative Solver

$*d = b$

$a = *d$



What is going to be
the next action in this
algorithm?



let $G = (V, E)$

$W = V$

while $W \neq []$ **do**

$n = \text{hd}(W)$

for each $v \in P(n)$ **do**

for each load " $a = *n$ " **do**

if $(v, a) \notin E$ **then**

$E = E \cup \{(v, a)\}$

$W = v::W$

for each store " $*n = b$ " **do**

if $(b, v) \notin E$ **then**

$E = E \cup \{(b, v)\}$

$W = b::W$

for each $(n, z) \in E$ **do**

$P(z) = P(z) \cup P(n)$

if $P(z)$ has changed **then**

$W = z::W$

Fernando Magno Quintão Pereira

`fernando@dcc.ufmg.br`

`lac.dcc.ufmg.br`

COLLAPSING CYCLES



The Problems with the Iterative Solver









- Finding the transitive closure of a graph has an $O(n^3)$ algorithm.
 - This algorithm works well for small graphs, but it is very slow for very large programs.
 - Programs like the kernel of Linux.
- But, in the early 2000's, researchers have made a very interesting observation:
 - All the nodes in a cycle always have the same points-to set.

1) Why is this observation true?

2) How can we use it to improve our algorithm?

Collapsing Cycles

- The key idea to capitalize on cycles is to collapse them, once we find them.



How to identify
cycles in the
graph?

Cycle Identification

- Collapsing cycles is very important: it is the difference between an algorithm that can be used, and one that cannot be used in practice.
- Although identifying cycles in graphs is easy, e.g., we can find them with a depth-first traversal of the graph, using this strategy in practice is not trivial.

1) What is the problem with using a DFS to find cycles?

2) But how can we still capitalize on cycles?

Lazy Cycle Detection

- Compiler researchers have identified many ways to detect cycles during the construction of the transitive closure:
 - Wave propagation
 - Deep propagation
 - Lazy cycle detection
- We shall be describing the lazy cycle detection approach♣, because it is easy to write in pseudo-code.

♣: The ant and the grasshopper: fast and accurate pointer analysis for millions of lines of code, 2007



Lazy Cycle Detection

- If we connect two nodes v_1 and v_2 that have the same points-to set, chances are that we are in the middle of a cycle.
 - We can try to find a cycle starting from v_2
 - If we can indeed find a cycle, then we collapse it
 - Otherwise we mark the edge (v_1, v_2) to avoid future searches that are likely to fail.

```
let  $G = (V, E)$   
 $R = \{\}$   
 $W = V$   
while  $W \neq []$  do  
   $n = \text{hd}(W)$   
  for each  $v \in P(n)$  do  
    for each load " $a = *n$ " do  
      if  $(v, a) \notin E$  then  
         $E = E \cup \{(v, a)\}$   
         $W = v :: W$   
    for each store " $*n = b$ " do  
      if  $(b, v) \notin E$  then  
         $E = E \cup \{(b, v)\}$   
         $W = b :: W$   
  for each  $(n, z) \in E$  do  
    if  $P(z) = P(n)$  and  $(n, z) \notin R$  then  
      DETECTANDCOLLAPSECYCLES( $z$ )  
       $R = R \cup \{(n, z)\}$   
   $P(z) = P(z) \cup P(n)$   
  if  $P(z)$  has changed then  
     $W = z :: W$ 
```

Lazy Cycle Detection

- We keep track of all the edges that have fired a cycle detection.
- In this way, we make sure that we do not try to find the same cycle over and over again, due to edges with the same points-to set.

```

let G = (V, E)
R = {}
W = V
while W ≠ [] do
  n = hd(W)
  for each v ∈ P(n) do
    for each load "a = *n" do
      if (v, a) ∉ E then
        E = E ∪ {(v, a)}
        W = v::W
    for each store "*n = b" do
      if (b, v) ∉ E then
        E = E ∪ {(b, v)}
        W = b::W
  for each (n, z) ∈ E do
    if P(z) = P(n) and (n, z) ∉ R then
      DETECTANDCOLLAPSECYCLES(z)
      R = R ∪ {(n, z)}
    P(z) = P(z) ∪ P(n)
    if P(z) has changed then
      W = z::W
  
```


Example of Lazy Cycle Detection

$c = \&d$

$e = \&a$

$a = b$

$b = c$

$c = *e$

How is the initial
points-to graph
for the program
on the left?

Example of Lazy Cycle Detection

$c = \&d$

$e = \&a$

$a = b$

$b = c$

$c = *e$

What will happen
once we process this
load constraint?


$\{a\}$
e

$\{d\}$ a

d

b
 $\{d\}$

c
 $\{d\}$



Example of Lazy Cycle Detection

$c = \&d$

$e = \&a$

$a = b$

$b = c$

$c = *e$

1) Does the inclusion of
this new edge trigger
cycle detection?

$\{a\}$
e


$\{d\}$ a

d

2) What do we do once
we find a cycle in the
constraint graph?

b
 $\{d\}$

c
 $\{d\}$



Example of Lazy Cycle Detection

$c = \&d$

$e = \&a$

$a = b$

$b = c$

$c = *e$

$\{a\}$
e

$\{d\}$ a/b/c


d

What are the advantages and the disadvantages of lazy cycle detection?

After we find a cycle, we collapse it into a single node. In this way, we reduce the quantity of nodes we must handle in every search and update operation of our algorithm.

Tradeoffs

- The main advantage of lazy cycle detection is the fact that it only fires a search for a cycle once the likelihood to find the cycle is high
- Another advantage is that lazy cycle detection is conceptually simple, and it is easy to implement
- The main drawback is that a cycle may remain on a points-to graph for a while until we detect it
- Another problem is that many searches still fail
 - The exact proportion of searches that fail has not been reported in the literature as of today.



Wave Propagation

- Wave propagation is another algorithm used to solve points-to analysis
- It relies on a fast technique[♠] to find strongly connected components in the constraint graph
- Once these components are found and collapsed, the algorithm propagates points to facts in waves, following a topological ordering of the newly modified graph
- This pattern continues, until the points-to graph stops changing.



Could you write
this descriptions
as pseudo-code?

[♠]: On Finding the Strongly Connected Components in a Directed Graph (1994)

Wave Propagation

```
repeat  
  changed = false  
  collapse Strongly Connected Components  
  WAVEPROPAGATION  
  ADDNEWEDGES  
  if a new edge has been added to G then  
    changed = true  
until changed = false
```

1) What do you think
WAVEPROPAGATION
does?

2) What about the
ADDNEWEDGES
algorithm?

Wave Propagation of Points-to Information

WAVEPROPAGATION(G, P, T)

while $T \neq []$

$v = \text{hd}(T)$

$P_{\text{dif}} = P_{\text{cur}}(v) - P_{\text{old}}(v)$

$P_{\text{old}}(v) = P_{\text{cur}}(v)$

if $P_{\text{dif}} \neq \{\}$

for each w such that $(v, w) \in E$ **do**

$P_{\text{cur}}(w) = P_{\text{cur}}(w) \cup P_{\text{dif}}$

Each node is associated with two sets, P_{cur} and P_{old} . The first denotes the current points-to facts that we know about the node. The second denotes the points-to fact that we knew in the last iteration of our algorithm.

The parameters of the algorithm are:

- G : the points-to graph
- P : the points-to facts
- T : the topological ordering of the nodes in G

The Creation of New Edges

ADDNEWEDGES($G = (E, V)$, C)

for each operation c such as $l = *r \in C$ **do**

$$P_{\text{new}} = P_{\text{cur}}(r) - P_{\text{cache}}(c)$$

$$P_{\text{cache}}(c) = P_{\text{cache}}(c) \cup P_{\text{new}}$$

for each $v \in P_{\text{new}}$ **do**

if $(v, l) \notin E$ **then**

$$E = E \cup \{(v, l)\}$$

$$P_{\text{cur}}(l) = P_{\text{cur}}(l) \cup P_{\text{old}}(v)$$

for each operation c such as $*l = r$ **do**

$$P_{\text{new}} = P_{\text{cur}}(l) - P_{\text{cache}}(c)$$

$$P_{\text{cache}}(c) = P_{\text{cache}}(c) \cup P_{\text{new}}$$

for each $v \in P_{\text{new}}$ **do**

if $(r, v) \notin E$ **then**

$$E = E \cup \{(r, v)\}$$

$$P_{\text{cur}}(v) = P_{\text{cur}}(v) \cup P_{\text{old}}(r)$$

We keep track of $P_{\text{cache}}(c)$, the last collection of points used in the evaluation of the complex constraint c . This optimization reduces the number of edges that must be checked for inclusion in G . $P_{\text{cache}}(c)$ is initially set to $\{\}$.

These updates will set up the ground for the next iteration of the wave propagation.

Wave Propagation: an example

- Let's illustrate the wave propagation algorithm with the following set of statements:

`h = &c`

`e = &g`

`b = c`

`h = &g`

`h = a`

`c = b`

`a = &e`

`f = d`

`b = a`

`d = *h`

`*e = f`

`f = &a`

What is the initial points-to graph for this program?

Answer explaining how many nodes, and how many edges this graph will have.

Wave Propagation: an example

$h = \&c$

$e = \&g$

$b = c$

$d = *h$

$h = \&g$

$h = a$

$c = b$



$f = \&a$

$a = \&e$

$f = d$

$b = a$

$*e = f$



How will be the
graph after we
merge SCCs?

Wave Propagation: an example

$h = \&c$

$e = \&g$

$b = c$

$d = *h$

$h = \&g$

$h = a$

$c = b$



$f = \&a$

$a = \&e$

$f = d$

$b = a$

$*e = f$



How will be the graph after the first wave propagation?

Wave Propagation: an example

$h = \&c$

$e = \&g$

$b = c$

$d = *h$

$h = \&g$

$h = a$

$c = b$



$f = \&a$

$a = \&e$

$f = d$

$b = a$

$*e = f$



How will be the graph after we add the new edges? Which constraints do we have to consider now?

Wave Propagation: an example

$h = \&c$

$e = \&g$

$b = c$

$d = *h$

$h = \&g$

$h = a$

$c = b$


$f = \&a$

$a = \&e$

$f = d$


$b = a$

$*e = f$




Can you finish
the iterations of
the algorithm in
this graph?

Wave Propagation: an example





What is the complexity of the wave propagation solver?
What about the lazy cycle detector?



After we collapse SCCs, the sequence of wave propagations causes no more changes, and no further edge is added to the graph. We have reached a fixed point, and we can stop the algorithm.

The Shape of the Heap

- From the result of the points-to analysis we can infer the shape that the heap may assume after the program executes.



STEENSGAARD'S ANALYSIS



DCC 888

Equality vs Subset Inclusion

- In the previous approach to pointer analysis, e.g., Andersen Style, a statement like $p = q$ means that everything that is pointed by q could be also pointed by p .

$$P(p) \supseteq P(q)$$

- We can speed up this analysis, by unifying both sides, instead of using the subset relation.

$$P(p) = P(q)$$

- This way of doing pointer analysis, e.g., based on the unification of both sides of an assignment, is called Steensgaard's analysis, after the work of *Bjarne Steensgaard*[♠].

Steensgaard's Analysis at Work

- We shall demonstrate how the unification based analysis works with an example:


```
x = &a;  
y = &b;  
p = &x;  
p = &y;
```

What is the shape of the heap, as reported by Andersen's analysis, after this program runs?

Steensgaard's Analysis at Work

- We shall demonstrate how the unification based analysis works with an example:

```
x = &a;  
y = &b;  
p = &x;  
p = &y;
```



- Steensgaard's analysis proceed by joining into equivalence classes elements that appear together in assignments.
- An assignment like `x = &a` means that `'*x'` and `'a'` are in the same equivalence class, which is pointed by `x's` class.


Steensgaard's Analysis at Work

`x = &a;`

`y = &b;`

`p = &x;`

`p = &y;`



What will be the effect of **this** assignment?

An assignment like `x = &a` means that `*x` and `a` are in the same equivalence class, which is pointed by `x`'s class.



Steensgaard's Analysis at Work

```
x = &a;
```

```
y = &b;
```

```
p = &x;
```

```
p = &y;
```



What will be the effect of **the third** assignment?


Steensgaard's Analysis at Work

```
x = &a;
```

```
y = &b;
```

```
p = &x;
```

```
p = &y;
```



And what will
happen after the
last statement is
analyzed?


Steensgaard's Analysis at Work

```
x = &a;
```

```
y = &b;
```

```
p = &x;
```

```
p = &y;
```



How fast can we
perform **this**
unification?

After the assignment $p = \&y$, we
have that $*p$ appears in two different
classes. We must unify them.


Steensgaard's Analysis at Work

`x = &a;`


`y = &b;`

`p = &x;`

`p = &y;`




But the process does not stop in this unification of the classes of `x` and `y`. It must be propagated throughout the children of these two classes, until we have only one line of pointers.



Union-Find[♡]

- This process of chain unifications has a fast implementation.
 - It can be implemented to run in $\alpha(n)$, where n is the number of elements to be unified, and α is the inverse Ackermann's function.
- The algorithm that implements this unification is called **union-find**.


By the way, why are these chains of unifications necessary?




Why are Chains of Unification Necessary?

What is the shape of the heap that we have inferred with this analysis?


Imagine that we had an assignment like $*p = q$. We need to find everything that is pointed by p , to unify it with q . If we had kept the equivalence classes of a and b separated, we would have to update them separately as well. Because they are the same, we can update them together, in a single pass.



Steensgaard's is less precise than Andersen's



Steensgaard's analysis is very fast; however, this speed pays a price in terms of precision. In this case, we have found out that x can point to either a or b . Same thing for y . **Andersen's** analysis would tell us that x can only point to a , and y can only point to b .



A Common Pattern


- All the algorithms that we have seen so far, follow a very common pattern: iterate until a fixed point is reached.
 - If no changes have been found in one iteration, no more changes will ever happen.
- This pattern happens in several different analyses:
 - All the data-flow analyses
 - The control flow analysis
 - The points-to analysis

Can you provide some intuition on why the different points-to solvers are guaranteed to terminate?

Flow Sensitiveness

- Even Andersen's analysis is not very precise, as it is flow insensitive.

1) `x = &a;`
2) `y = &b;`
3) `p = &x;`
4) `p = &y;`



The points-to graph represents any possible edge that can exist at any moment during the execution of the program. Some of these edges, of course, cannot exist at every program point. Some of them cannot even exist at the same time.

Which edges can exist after instruction 1 executes? What about after instructions 2, 3 and 4?

Flow Sensitiveness

What is the complexity of the size of the output of shape analysis?





1) $x = \&a;$

2) $y = \&b;$

3) $p = \&x;$

4) $p = \&y;$

There are analyses that are strong enough to track the shape of the points-to graph at each program point. One such analysis is "Shape analysis"[♡]. You can imagine that it is pretty costly...



♡: Parametric shape analysis via 3-valued logic, 2002

A Bit of History

- Inclusion-based points-to analysis has been described by Lars Andersen in his PhD dissertation.
- Unification-based points-to analysis was an idea due to Bjarne Steensgaard, in the mid 90's.
- Lazy cycle detection was invented by Ben Hardekopf in 2007.
- Wave propagation was designed by Pereira and Berlin in 2011.

- Andersen, L. "Program Analysis and Specialization for the C Programming Language", PhD Thesis, University of Copenhagen, (1994)
- Hardekopf, B. and Lin, C. "The Ant and the Grasshopper: fast and accurate pointer analysis for millions of lines of code", PLDI, pp 290-299 (2007)
- Pereira, F. and Berlin, D. "Wave Propagation and Deep Propagation for Pointer Analysis", CGO, pp 126-135 (2009)
- Steensgaard, B., "Points-to Analysis in Almost Linear Time", POPL, pp 32-41 (1995)